



yugabyteDB

# Migrating MySQL to Yugabyte using `ysql_loader`

Mike Lee - Solutions Architect

Suranjan Kumar - Ecosystem Integration Engineer

Yugabyte


# Agenda


---

- What is `ysql_loader`?
- Where can I Git it?
- How to use it?
  - configuration options
- How to run it!
- Yugabyte features added
- Demo

# What is ysql\_loader

- ysql\_loader = pgloader (pgloader.io)
- Dimitri Fontaine: wrote and maintains of pgloader, Major Contributor to PostgreSQL, author of The Art of PostgreSQL
  - Open Source
  - ysql\_loader forked from pgloader (Suranjan)





## PGLOADER

[BLOG](#) [ABOUT](#) [LICENSING](#) [ROADMAP](#) [SERVICES](#) [WHITE PAPER](#)

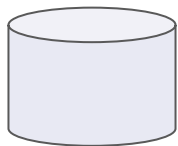
pgloader loads data into PostgreSQL and allows you to implement [Continuous Migration](#) from your current database to PostgreSQL. Read the [White Paper](#) to learn how to limit risks and control your budget, and start your PostgreSQL migration today!

# ysql\_loader

---

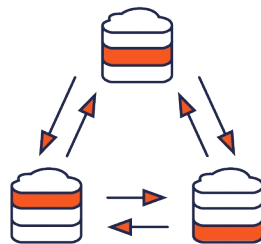
- Migrates database objects (tables, indexes, sequences) and loads data from files (CSV, Fixed, DBF, IXF) or directly connecting to source databases such as MySQL, SQL Lite, MS SQL and PostgreSQL
- Can load from flat files (CSV, Fixed Format) or directly connect and migrate entire databases to
- Follows a typical database migration workflow:
  - Create the target database on Yugabyte
  - Gather database objects from source database metadata catalog
  - Create Tables and Indexes
  - Copy table data using PostgreSQL COPY
  - Add the constraints, primary & foreign keys, and comments
- Has a rich set of options that allow you to customize ysql\_loader to your exact needs.

# Where to install ysql\_loader?



MySQL, MS SQL,  
SQL Lite, Postgres

ysql\_loader

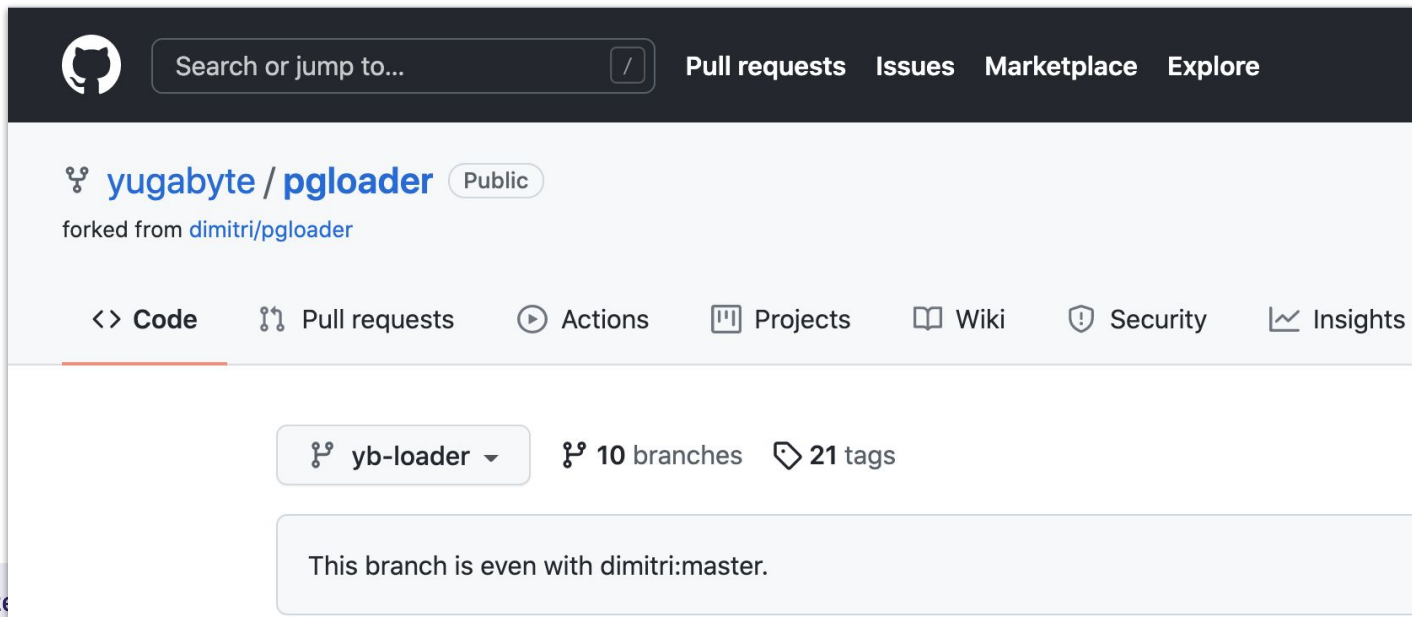


For production migrations, ysql\_loader should be installed on a separate server.  
For this workshop, ysql\_loader can be installed on the same server as MySQL

# Pgloader - how to Git it...

\$git clone <https://github.com/yugabyte/pgloader>

NOTE: install pgloader on a separate “migration server”



## pgloader - build it and they will migrate....

---

```
$ git clone https://github.com/yugabyte/pgloader
```

```
#
```

### # GET ALL LIBRARIES

```
$ apt-get install sbcl unzip libsqlite3-dev make curl gawk freetds-dev libzip-dev
```

```
$ cd /path/to/pgloader
```

```
$ make pgloader
```

```
#
```

### # purposely renamed, but you could replace pgloader

```
$ sudo cp /build/bin/pgloader /usr/bin/ysql_loader
```

## pgloader(Docker) - build it and they will migrate

---

```
$ git clone https://github.com/yugabyte/pgloader
```

```
$ cd /path/to/pgloader
```

```
$ sudo docker -t ysql-loader:v1.3 build .
```



# pgloader command-file options

---

## Configuration Sections

- **Database Source: FROM**
- **Migration Options: WITH**
- **Casting Rules: CAST**
- **Partial Migration: INCLUDING/EXCLUDING NAMES**
- **Encoding Support**
- **Schema Transformations**
- **View Support**

Check out: <https://pgloader.readthedocs.io/en/latest/index.html>

# Pgloader command options - FROM

---

## LOAD DATABASE

```
FROM mysql://root:<password>@IP.Addr:3306/ml_migratedb  
INTO postgresql://yugabyte:<password>@IP.Addr:5433/ml_migratedb;
```

Determines the source and targets

Pre-migration checks: test the login and password and IP to make sure you can access both the source database and YugabyteDB from the wherever ysql\_loader was installed.

# Pgloader command options - WITH

```
LOAD DATABASE
```

```
FROM mysql://root:<password>@IP.Ad.dre.ss:3306/ml_migratedb
```

```
INTO postgresql://yugabyte:<password>@IP.Ad.dres.ss:5433/ml_migratedb
```

```
WITH batch rows = 1000, truncate
```

```
;
```

Default WITH options with MySQL:

- no truncate
- create tables
- include drop
- create indexes
- reset sequences
- foreign keys
- **downcase identifiers** - be mindful of tables with same name but differ by capitalization.
- **uniquify index names** - PG index names have to be unique per-schema (MySQL is per-table)

# Pgloader command options - INCLUDING/EXCLUDING

---

LOAD DATABASE

FROM mysql://root:<password>@IP.Ad.dres.ss:3306/ml\_migratedb

INTO postgresql://yugabyte:<password>@IP.Ad.dres.ss:5433/ml\_migratedb

WITH batch rows = 1000, truncate

INCLUDING ONLY TABLE NAMES MATCHING ~/ml\_/, 'orders\_orig'

EXCLUDING TABLE NAMES MATCHING ~<orig>

;

INCLUDING: comma separated list of table names or regular expression used to limit the tables.

EXCLUDING: comma separated list of table names or regular expressions used to limit the tables, HOWEVER, This filter only affects the result of the INCLUDING filter.

# Pgloader command options - CAST

```
LOAD DATABASE
```

```
FROM mysql://root:<password>@IP.Ad.dres.ss:3306/ml_migratedb
```

```
INTO postgresql://yugabyte:<password>@IP.Ad.dres.ss:5433/ml_migratedb
```

```
WITH batch rows = 1000, truncate
```

```
INCLUDING ONLY TABLE NAMES MATCHING ~/ml_/, 'orders_orig'
```

```
EXCLUDING TABLE NAMES MATCHING ~<orig>
```

```
CAST
```

```
type tinyint to smallint drop typemod
```

```
;
```

MySQL TINYINT(1) supports values 0,1,2,3

PostgreSQL has BOOLEAN, but it only supports

- TRUE, 't', 'true', 'y', 'yes', 'on', '1' or FALSE, 'f', 'false', 'n', 'no', 'off', '0'
- To retain values 2 and 3, you must convert the data type to a SMALLINT.

CAST does this automatically for you!

# mysql\_loader command options - CAST

## MYSQL

```
mysql> show full columns from ml_tinyint1;
```

| Field  | Type                      | Collation         | Null | Key | Default | Extra          |
|--------|---------------------------|-------------------|------|-----|---------|----------------|
| id     | int(10) unsigned zerofill | NULL              | NO   | PRI | NULL    | auto_increment |
| ti_col | tinyint(1)                | NULL              | YES  |     | NULL    |                |
| name   | varchar(20)               | latin1_swedish_ci | NO   |     | NULL    |                |

## YugabyteDB

Converted to SMALLINT

```
ml_migratedb=# \d+ ml_tinyint1
```

| Column | Type                  | Collation | Nullable | Default                                 |
|--------|-----------------------|-----------|----------|---|
| id     | bigint                |           | not null | nextval('ml_tinyint1_id_seq'::regclass) |
| ti_col | smallint              |           |          |   |
| name   | character varying(20) |           | not null |   |

Indexes:  
"ml\_tinyint1\_pkey" PRIMARY KEY, lsm (id HASH)

# Data type considerations

---

| Data Type       | MySQL  | PostgreSQL                            |
|-----------------|--|---------------------------------------|
| Integers(bytes) | TINYINT(1), SMALLINT(2),<br>MEDIUMINT(3), INT(4),<br>BIGINT(8) | SMALLINT(2), INTEGER(4),<br>BIGINT(8) |
| Text            | TINYTEXT, TEXT,<br>MEDIUMTEXT, LONGTEXT                        | TEXT                                  |
| Number          | DECIMAL, NUMERIC   | DECIMAL, NUMERIC                      |
| Double          | DOUBLE   | DOUBLE                                |
|                 |  |                                       |

# mysql\_loader command options - RESET SEQUENCES

MySQL

```
mysql> select max(id) from ml_order_line;
```

```
+-----+  
| max(id) |  
+-----+  
|  924280 |  
+-----+
```

MySQL: ml\_order\_line table has a sequence called **ml\_order\_line\_id\_seq**

Yugabyte

In ybdb: the max value is preserved and next value ready to use.

```
[ml_migratedb=# select nextval('ml_order_line_id_seq');  
nextval  
-----  
 924281  
(1 row)
```



# mysql\_loader command options - DUMPDDL ONLY

## LOAD DATABASE

FROM mysql://root:P8ssw0rd2@10.142.0.2:3306/ml\_migratedb

INTO postgresql://yugabyte:yugabyte@10.204.0.5:5433/ml\_migratedb

## WITH dumpddl only

including only table names matching ~/ml\_/, 'orders\_orig'

## CAST

type tinyint to smallint drop typemod

;

DDL extracted to file ddl.sql !

| 2021-09-17T22:23:22.396000Z LOG report summary reset |                        |        |      |          |       |            |      |       |  |
|--|------------------------|--------|------|----------|-------|------------|------|-------|--|
|  | table name             | errors | read | imported | bytes | total time | read | write |  |
|  | fetch meta data        | 0      | 6    | 6        | 0     | 0.095s     |      |       |  |
|  | Create Schemas         | 0      | 0    | 0        | 0     | 0.007s     |      |       |  |
|  | Create SQL Types       | 0      | 0    | 0        | 0     | 0.007s     |      |       |  |
|  | Create tables          | 0      | 8    | 8        | 0     | 17.277s    |      |       |  |
|  | Set Table OIDs         | 0      | 4    | 4        | 0     | 0.012s     |      |       |  |
|  | create primary key     | 0      | 3    | 3        | 51000 | 16.375s    |      |       |  |
|  | Index Build Completion | 0      | 2    | 2        | 0     | 10.008s    |      |       |  |
|  | Create Indexes         | 0      | 2    | 2        | 0     | 8.731s     |      |       |  |
|  | Reset Sequences        | 0      | 0    | 0        | 0     | 1.076s     |      |       |  |
|  | Primary Keys           | 0      | 0    | 0        | 0     | 0.000s     |      |       |  |
|  | Create Foreign Keys    | 0      | 0    | 0        | 0     | 0.000s     |      |       |  |
|  | Create Triggers        | 0      | 0    | 0        | 0     | 0.001s     |      |       |  |
|  | Install Comments       | 0      | 1    | 1        | 0     | 0.026s     |      |       |  |
|  | Total import time      | ✓      | 0    | 0        | 0     | 19.842s    |      |       |  |

# mysql\_loader command options - DATA ONLY

## LOAD DATABASE

FROM mysql://root:P8ssw0rd2@10.142.0.2:3306/ml\_migratedb

INTO postgresql://yugabyte:yugabyte@10.204.0.5:5433/ml\_migratedb

## WITH data only

including only table names matching ~/ml\_/, 'orders\_orig'

## CAST

type tinyint to smallint drop typemod

;

| 2021-09-17T22:54:10.801000Z | LOG    | report | summary  | reset   |            |           |           |  |  |
|-----------------------------|--------|--------|----------|---------|------------|-----------|-----------|--|--|
| table name                  | errors | read   | imported | bytes   | total time | read      | write     |  |  |
| fetch meta data             | 0      | 4      | 4        |         | 0.086s     |           |           |  |  |
| Drop Foreign Keys           | 0      | 0      | 0        |         | 0.000s     |           |           |  |  |
| ml_migratedb.ml_order_line  | 0      | 800000 | 800000   | 56.9 MB | 2m32.413s  | 2m11.162s | 2m29.125s |  |  |
| ml_migratedb.ml_orders      | 0      | 80000  | 80000    | 3.5 MB  | 21.583s    | 0.632s    | 17.456s   |  |  |
| ml_migratedb.ml_tinyint1    | 0      | 201    | 201      | 4.3 kB  | 21.906s    | 0.002s    |           |  |  |
| ml_migratedb.orders_orig    | 0      | 30000  | 30000    | 1.2 MB  | 3.024s     | 0.214s    | 0.121s    |  |  |
| COPY Threads Completion     | 0      | 4      | 4        |         | 2m32.395s  |           |           |  |  |
| Reset Sequences             | 0      | 0      | 0        |         | 2.063s     |           |           |  |  |
| Create Foreign Keys         | 0      | 0      | 0        |         | 0.000s     |           |           |  |  |
| Install Comments            | 0      | 1      | 1        |         | 0.025s     |           |           |  |  |
| Total import time           | ✓      | 910201 | 910201   | 61.6 MB | 2m34.483s  |           |           |  |  |

# Running pgloader

pgloader can be run by passing arguments or by passing a command-file

```
ybploader --help
ybploader [ option ... ] command-file ...
ybploader [ option ... ] SOURCE TARGET

--help -h          boolean Show usage and exit.
--version -V       boolean Displays pgloader version and exit.
--quiet -q         boolean Be quiet
--verbose -v       boolean Be verbose
--debug -d         boolean Display debug level information.
--client-min-messages string Filter logs seen at the console (default:
"warning")
--log-min-messages string Filter logs seen in the logfile (default:
"notice")
--summary -S       string Filename where to copy the summary
--root-dir -D      string Output root directory. (default:
#P"/tmp/pgloader/")
--upgrade-config -U boolean Output the command(s) corresponding
to .conf file for v2.x
--list-encodings -E boolean List pgloader known encodings and exit.
--logfile -L       string Filename where to send the logs.
--load-lisp-file -l string Read user code from files
```

```
--dry-run          boolean Only check database connections, don't
load anything.
--on-error-stop    boolean Refrain from handling errors
properly.
--no-ssl-cert-verification boolean Instruct OpenSSL to bypass
verifying certificates.
--context -C       string Command Context Variables
--with            string Load options
--set             string PostgreSQL options
--field          string Source file fields specification
--cast           string Specific cast rules
--type           string Force input source type
--encoding       string Source expected encoding
--before        string SQL script to run before loading the data
--after         string SQL script to run after loading the data
--self-upgrade   string Path to pgloader newer sources
--regress       boolean Drive regression testing
```

# Running pgloader

---

Using the command file:

```
sudo ysql_loader --verbose <cmd_file>
```

Or by running it from Docker...

```
docker run --rm --name pgloader \  
  -v <local_dir pgloader_config_dir>:<mount_path_in_container> \  
  yugabytedb/pgloader:v1.1 pgloader --verbose \  
  <mount_path_in_container>/<pgloader_config_file_in_local_dir>
```

example:

```
sudo docker run --rm --name pgloader \  
  -v /home/ubuntu/Mlee/PgloaderDSS:/tmp \  
  /bin/bash -c 'cd /tmp; pgloader --verbose \  
  /tmp/schmonly.conf'
```



**Let's migrate some data!**

# Migration output using --verbose option

```
2021-09-17T22:07:33.019000Z NOTICE Starting pgloader, log system is ready.
2021-09-17T22:07:33.046000Z LOG pgloader version "3.6.15edf6f"
2021-09-17T22:07:33.322000Z LOG Migrating from #<MYSQL-CONNECTION mysql://stan@10.142.0.2:3306/ml_migratedb {10086B2593}>
2021-09-17T22:07:33.322000Z LOG Migrating into #<PGSQL-CONNECTION pgsql://yugabyte@10.204.0.5:5433/ml_migratedb {10086B3D23}>
>
2021-09-17T22:07:33.473000Z NOTICE Prepare PostgreSQL database.
2021-09-17T22:07:52.187000Z NOTICE Processing tables in this order: ml_migratedb.ml_order_line: 807731 rows, ml_migratedb.ml_orders: 80254 rows, ml_migratedb.orders_orig: 28379 rows, ml_migratedb.ml_tinyint1: 202 rows
2021-09-17T22:07:52.187000Z NOTICE Executing SQL block for create primary key and sql is (ALTER TABLE ml_migratedb.ml_order_line ADD PRIMARY KEY (id, ol_w_id, ol_d_id, ol_o_id, ol_number);)
2021-09-17T22:07:52.793000Z NOTICE executing command: ALTER TABLE ml_migratedb.ml_order_line ADD PRIMARY KEY (id, ol_w_id, ol_d_id, ol_o_id, ol_number);
2021-09-17T22:07:57.675000Z NOTICE Executing SQL block for create primary key and sql is (ALTER TABLE ml_migratedb.ml_orders ADD PRIMARY KEY (id, o_w_id, o_d_id, o_id);)
2021-09-17T22:07:58.293000Z NOTICE executing command: ALTER TABLE ml_migratedb.ml_orders ADD PRIMARY KEY (id, o_w_id, o_d_id, o_id);
2021-09-17T22:08:03.208000Z NOTICE Executing SQL block for create primary key and sql is (ALTER TABLE ml_migratedb.ml_tinyint1 ADD PRIMARY KEY (id);)
2021-09-17T22:08:03.822000Z NOTICE executing command: ALTER TABLE ml_migratedb.ml_tinyint1 ADD PRIMARY KEY (id);
2021-09-17T22:08:08.604000Z NOTICE Done with prepare postgres and executing after-schema
2021-09-17T22:08:08.605000Z NOTICE Processing tables in this order: ml_migratedb.ml_order_line: 807731 rows, ml_migratedb.ml_orders: 80254 rows, ml_migratedb.orders_orig: 28379 rows, ml_migratedb.ml_tinyint1: 202 rows
2021-09-17T22:08:08.619000Z NOTICE COPY ml_migratedb.ml_order_line with 807731 rows estimated [0/4]
2021-09-17T22:08:08.629000Z NOTICE COPY ml_migratedb.ml_orders with 80254 rows estimated [3/4]
2021-09-17T22:08:09.285000Z NOTICE COPY ml_migratedb.orders_orig with 28379 rows estimated [3/4]
2021-09-17T22:08:19.779000Z NOTICE COPY ml_migratedb.ml_tinyint1 with 202 rows estimated [2/4]
2021-09-17T22:09:19.204000Z NOTICE Executing CREATE UNIQUE INDEX idx_18014_order_line_i1 ON ml_migratedb.ml_order_line (id);
2021-09-17T22:10:13.955000Z NOTICE Executing CREATE UNIQUE INDEX idx_18008_orders_i1 ON ml_migratedb.ml_orders (id);
2021-09-17T22:10:22.443000Z NOTICE Completing PostgreSQL database.
2021-09-17T22:10:22.443000Z NOTICE Reset sequences
```

# The migrated database summary

| 2021-09-17T22:10:25.494000Z LOG report summary reset |        |        |          |                  |           |            |         |          |
|--|--------|--------|----------|------------------|-----------|------------|---------|----------|
| table name   | errors | read   | imported | # yb-hw          | bytes     | total time | read    | write    |
| -----  | -----  | -----  | -----    | # yb-k8s         | -----     | -----      | -----   | -----    |
| fetch meta data                                      | 0      | 6      | #6       | yb-platform      | 0.100s    |            |         |          |
| Create Schemas                                       | 0      | 0      | 0        | yb-releases      | 0.006s    |            |         |          |
| Create SQL Types                                     | 0      | 0      | 0        |                  | 0.008s    |            |         |          |
| Create tables  | 0      | 8      | #8       | yb-support-all   | 18.274s   |            |         |          |
| Set Table OIDs                                       | 0      | 4      | #4       | yb-support-amex  | 0.012s    |            |         |          |
| create primary key                                   | 0      | 3      | 3        | yb-support-cisco | 16.418s   |            |         |          |
| -----  | -----  | -----  | -----    | -----            | -----     | -----      | -----   | -----    |
| ml_migratedb.ml_order_line                           | 0      | 800000 | 800000   | yb-s             | 56.9 MB   | 1m10.368s  | 57.023s | 1m7.093s |
| ml_migratedb.ml_orders                               | 0      | 80000  | 80000    |                  | 3.5 MB    | ↓ 17.942s  | 0.666s  | 5.001s   |
| ml_migratedb.orders_orig                             | 0      | 30000  | 30000    |                  | 1.2 MB    | 10.493s    | 0.295s  | 0.078s   |
| ml_migratedb.ml_tinyint1                             | 0      | 201    | 201      |                  | 4.3 kB    | 0.038s     | 0.002s  | 0.001s   |
| -----  | -----  | -----  | -----    | -----            | -----     | -----      | -----   | -----    |
| COPY Threads Completion                              | 0      | 4      | 4        |                  | 1m10.339s |            |         |          |
| Index Build Completion                               | 0      | 2      | 2        |                  | 1m3.443s  |            |         |          |
| Create Indexes                                       | 0      | 2      | 2        |                  | 1m0.745s  |            |         |          |
| Reset Sequences                                      | 0      | 0      | 0        |                  | 2.386s    |            |         |          |
| Primary Keys   | 0      | 0      | 0        |                  | 0.000s    |            |         |          |
| Create Foreign Keys                                  | 0      | 0      | 0        |                  | 0.000s    |            |         |          |
| Create Triggers                                      | 0      | 0      | 0        |                  | 0.001s    |            |         |          |
| Install Comments                                     | 0      | 1      | 1        |                  | 0.025s    |            |         |          |
| -----  | -----  | -----  | -----    | -----            | -----     | -----      | -----   | -----    |
| Total import time                                    | ✓      | 910201 | 910201   |                  | 61.6 MB   | 3m16.939s  |         |          |

# Migrating MySQL to Yugabyte best practices

---

The workflow is similar, but we suggest the following extra steps:

- Gather row counts from all tables being migrated
- If > 40 tables, Yugabyte Universe should set `ysql_num_shards_per_tserver = 1`, otherwise use the default of `ysql_num_shards_per_tserver = 8`.
  - Any tables and indexes with > 1M rows need to be split into 24 tablets by adding `SPLIT INTO X TABLETS` on `CREATE TABLE` statement
    - `SPLIT INTO` clause will specify the number of tablets to be created for the table. The hash range is then evenly split across those tablets.
    - Any tables with < 1M rows will not need to be `SPLIT` and will automatically be created with 1 shard/tablet
- Use the `DATA ONLY` option to move data after database and tables have been created



# Yugabyte Branch of pgloader

---

- We create primary key on empty tables avoiding an expensive operation in Yugabyte
- We create indexes in single thread after table completion
- We have option to change the table schema to add YugabyteDB constructs like number of tablets, changing primary key
- We have avoided using savepoints and use Yugabyte's COPY command

## Summary

---

- Migrate MySQL, MS SQL, SQL Lite, and Postgres databases to YugabyteDB
- Flexible command options to customize the migration
- Proven, Open Source migration tool
- Optimized to work with YugabyteDB

# References/Acknowledgements

---

- YSQL Loader: <https://docs.yugabyte.com/latest/integrations/ysql-loader/#root>
- pgloader documentation: <https://pgloader.readthedocs.io/en/latest/intro.html>
- Yugabyte Branch of pgloader: <https://github.com/yugabyte/pgloader>



yugabyte**DB**

# Thank You

Join us on Slack: [yugabyte.com/slack](https://yugabyte.com/slack)

Star us on Github: [github.com/yugabyte/yugabyte-db](https://github.com/yugabyte/yugabyte-db)