

Migrating MySQL to YugabyteDB Using ysql_loader Workshop

V1.0 Original Author: Eric Pratt, Taylor Mull, Stanley Sung

V1.1 Mike Lee

We understand that database migrations can be painful. We have helped users successfully migrate from MySQL to YugabyteDB, a PostgreSQL-compatible distributed SQL database. A very popular tool to accomplish this task is [pgloader](#). In this post, we will cover how to migrate both your MySQL schema and data to YugabyteDB.

Prerequisites

Before starting the migration there are a few prerequisites you'll need to address for the Workshop.

Source Database:

You will need a “source” database to access. Supported databases are: MySQL, MS SQL, SQL Lite, or PostgreSQL. For this workshop we will be migrating from a MySQL database.

To get started creating a MySQL database use this link:

<https://dev.mysql.com/doc/mysql-getting-started/en/>

Yugabyte Database:

You can quickly install and create a single node Yugabyte cluster here:

<https://docs.yugabyte.com/latest/quick-start/>

Ysqlsh command line tool:

Download and install [ysqlsh](#) command line tool with connectivity to a running YugabyteDB cluster that you are going to migrate into.

Create a database on the YugabyteDB cluster that you will migrate into from the source cluster. This database needs to match the database name from MySQL. There would not be any tables or data behind the database but pgloader requires the database to exist on the target cluster. You can create the database by following these steps:

```
# Use ysqlsh to connect to any node of the YugabyteDB cluster
ysqlsh --host=<ip>
```

```
# Create the database
create database <name>;
```

Preparing to Migrate

For the workshop, we will install the `mysql_loader` on the same server as MySQL.

Verify Connectivity to Source MySQL DB

Next, ensure that the migration machine is able to connect to the MySQL server, because `pgloader` will need to communicate with MySQL over port 3306. Also make sure that the username/password used by `pgloader` and the IP address from which `pgloader` is connecting has permissions to connect to the MySQL source database.

1. To verify network connectivity you can use `telnet <MySQL_ip> 3306`
2. To grant the `pgloader` ip the permissions to access the database you can run the following command.

```
GRANT ALL PRIVILEGES ON *.* TO 'root'@'<pgloader_instance_ip>' WITH
GRANT OPTION;

flush PRIVILEGES;
```

3. To add a password for the ip/user combination, use the following command:

```
SET PASSWORD FOR 'root'@'<pgloader_instance_ip>' =  
PASSWORD('<password>');
```

If you see a *“failed to connect”* message like this after doing the above steps check the access with your MySQL DBA:

```
2021-04-22T17:33:33.232901Z ERROR #1=mysql: Failed to connect to  
#1# at "172.161.20.87" (port 3306) as user "root": MySQL Error  
[1045]: "Access denied for user  
'root'@'ip-172-161-27-195.us-east-2.compute.internal' (using  
password: YES)"  
2021-04-22T17:33:33.232990Z LOG report summary reset  
      table name      errors      rows      bytes      total time  
-----  
      fetch meta data          0          0          0          0.000s  
-----  
-----
```

Create a test database on MySQL

Verify Connectivity to Target YugabyteDB Cluster

The final check is to make sure the pgloader instance is able to reach the YugabyteDB cluster. For this we'll need to check to make sure the pgloader instance is able to communicate with one of the YugabyteDB nodes across port 5433.

1. To verify connectivity run `telnet <YugabyteDB_node_ip> 5433`

Installing pgloader (command line)

```
$ git clone https://github.com/yugabyte/pgloader  
#  
# GET ALL LIBRARIES  
$ apt-get install sbcl unzip libsqlite3-dev make curl gawk  
freetds-dev libzip-dev  
$ cd /path/to/pgloader  
$ make pgloader
```

```
##purposely renamed, but you could replace pgloader
$ sudo cp build/bin/pgloader /usr/bin/ysql_loader
```

For pgloader command flags:

```
ysql_loader --help
```

```
pgloader [ option ... ] command-file ...
pgloader [ option ... ] SOURCE TARGET
  --help -h                      boolean  Show usage and exit.
  --version -V                   boolean  Displays pgloader version and exit.
  --quiet -q                    boolean  Be quiet
  --verbose -v                   boolean  Be verbose
  --debug -d                    boolean  Display debug level information.
  --client-min-messages         string   Filter logs seen at the console (default:
"warning")
  --log-min-messages           string   Filter logs seen in the logfile (default:
"notice")
  --summary -S                  string   Filename where to copy the summary
  --root-dir -D                 string   Output root directory. (default:
#P"/tmp/pgloader/")
  --upgrade-config -U           boolean  Output the command(s) corresponding to
.conf file for v2.x
  --list-encodings -E           boolean  List pgloader known encodings and exit.
  --logfile -L                  string   Filename where to send the logs.
  --load-lisp-file -l           string   Read user code from files
  --dry-run                     boolean  Only check database connections, don't
load anything.
  --on-error-stop               boolean  Refrain from handling errors properly.
  --no-ssl-cert-verification    boolean  Instruct OpenSSL to bypass verifying
certificates.
  --context -C                  string   Command Context Variables
  --with                        string   Load options
  --set                         string   PostgreSQL options
  --field                       string   Source file fields specification
  --cast                        string   Specific cast rules
  --type                        string   Force input source type
  --encoding                    string   Source expected encoding
  --before                      string   SQL script to run before loading the data
  --after                       string   SQL script to run after loading the data
  --self-upgrade                string   Path to pgloader newer sources
  --regress                     boolean  Drive regression testing
```

Setting Up pgloader Command File

Instead of using the pgloader command line SOURCE TARGET with options, we use a command file to simplify the docker command.

```
load database
```

```
from mysql://root:password@<IP>:3306/testdb
into postgresql://yugabyte:yugabyte@<IP>:5433/testdb
WITH
    max parallel create index=1, batch rows = 1000;
```

Running pgloader

pgloader Schema and Data Migration

The docker image is based on centos7. We can “bash” into the container and run the pgloader command (/usr/local/bin/pgloader). In this doc, we store the pgloader command-file in “centos” home directory and use docker volume to map the configuration from centos home directory to docker container directory. The docker flag “--rm” will make sure the container is removed once the pgloader job is done.

```
docker run --rm --name <name_for_container> \
    -v <local_dir pgloader_config_dir>:<mount_path_in_container> \
    yugabytedb/pgloader:v1.1 pgloader \
    <mount_path_in_container>/<pgloader_config_file>
```

example:

```
sudo docker run --rm --name pgloader \
-v /home/ubuntu/Mlee/PgloaderDSS:/tmp \
ybloader:v1.2 pgloader \
/tmp/mysql_yb.cmd
```

Ex. pgloader command

```
[root@ip-172-161-27-195 centos]# pwd
/home/centos
[root@ip-172-161-27-195 centos]# ls
pgloader.conf
[root@ip-172-161-27-195 centos]# docker run --rm --name pgloader1
-v /home/centos:/tmp yugabytedb/pgloader:v1.1 pgloader -v -L
/tmp/pgloader.log /tmp/pgloader.conf
```

We should get the below output when the pgloader is running.

```
2021-04-22T18:49:00.000672Z LOG pgloader version "3.6.3~devel"
2021-04-22T18:49:00.264485Z LOG Migrating from #<MYSQL-CONNECTION
mysql://root@172.161.20.87:3306/testdb #x302001D3B50D>
```

```
2021-04-22T18:49:00.264662Z LOG Migrating into #<PGSQL-CONNECTION
pgsql://yugabyte@172.161.20.43:5433/testdb #x302001D3B3AD>
```

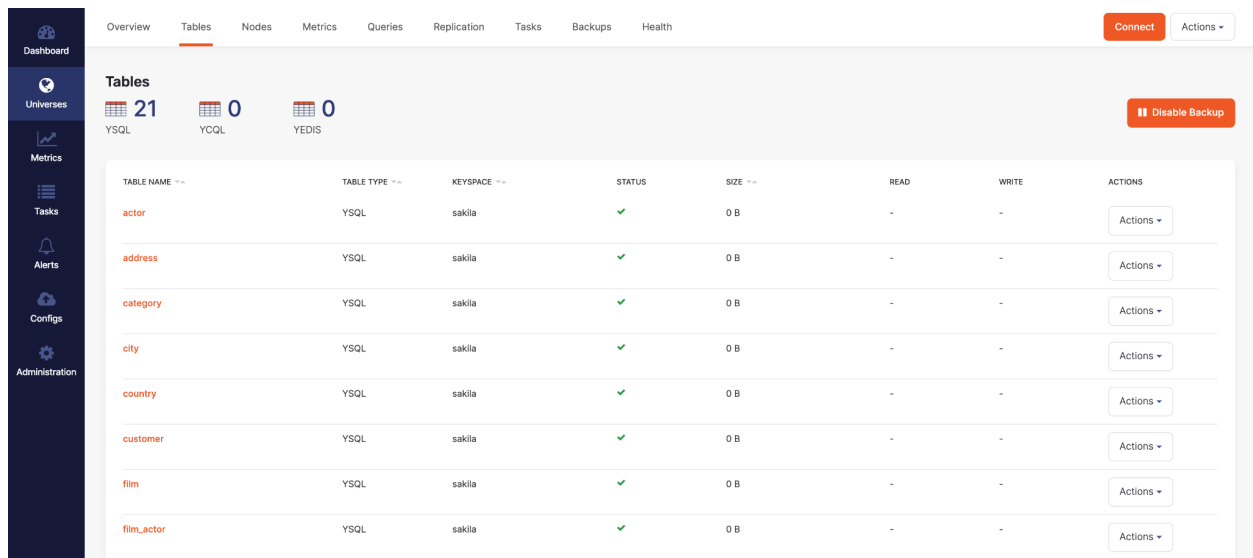
We can also double check the pgloader is running by running docker ps

```
[centos@ip-172-161-27-195 ~]$ sudo docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS
2dc611d4d412   yugabytedb/pgloader:v1.1           "pgloader /tmp/pgloa..." 16 minutes ago Up 16 minutes
pgloader1
```

We can also tail the pgloader.log we specified in the docker command.

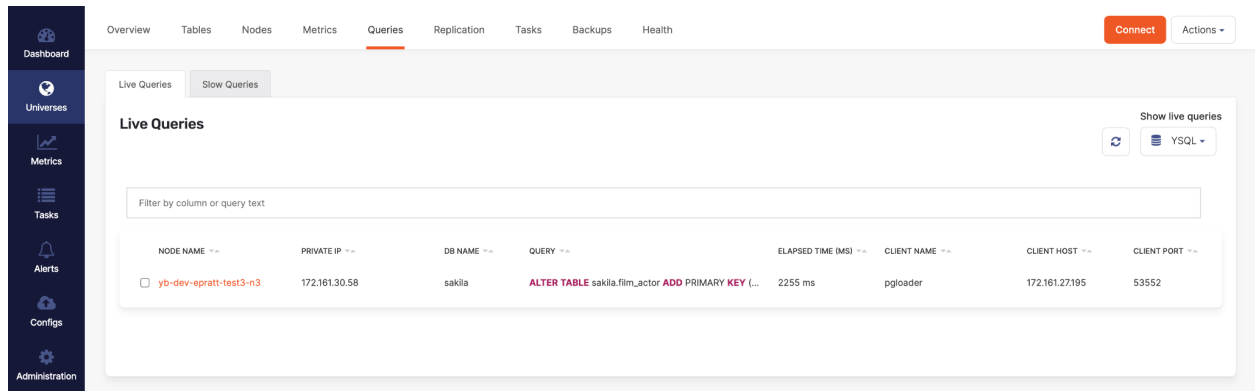
```
[centos@ip-172-161-27-195 ~]$tail -f pgloader.log
```

We can check where the pgloader is currently at by going to the platform UI and checking the tables section to see the tables start to load in.



Overview Tables Nodes Metrics Queries Replication Tasks Backups Health									
Connect Actions									
Tables									
21		0		0		Disable Backup			
TABLE NAME	TABLE TYPE	KEYSPACE	STATUS	SIZE	READ	WRITE	ACTIONS		
actor	YSQL	sakila	✓	0 B	-	-	Actions		
address	YSQL	sakila	✓	0 B	-	-	Actions		
category	YSQL	sakila	✓	0 B	-	-	Actions		
city	YSQL	sakila	✓	0 B	-	-	Actions		
country	YSQL	sakila	✓	0 B	-	-	Actions		
customer	YSQL	sakila	✓	0 B	-	-	Actions		
film	YSQL	sakila	✓	0 B	-	-	Actions		
film_actor	YSQL	sakila	✓	0 B	-	-	Actions		

In addition you can check the live queries and see current queries/ddl changes being made on the cluster at that time.



pgloader Schema Migration Only

If you only wanted to migrate the schema from MySQL to YugabyteDB and not include the data you can add `WITH schema only` to the pgloader command file and pgloader will only load the schema.

```
load database
from mysql://root:password@172.161.20.87:3306/testdb
into postgresql://yugabyte:yugabyte@172.161.30.169:5433/testdb
WITH
    max parallel create index=1, batch rows = 1000, schema only;
```

Validation

Once pgloader finishes the migration, you will get a summary of the migration steps which includes how long each step took and the number of rows inserted.

2021-09-21T19:42:05.153000Z LOG report summary reset								
table name	errors	read	imported	bytes	total time	read	write	
fetch meta data	0	4	4		0.105s			
Drop Foreign Keys	0	0	0		0.000s			
Truncate	0	4	4		0.535s			
ml_migratedb.ml_order_line	0	800000	800000	56.9 MB	59.317s	49.590s	56.979s	
ml_migratedb.ml_orders	0	80000	80000	3.5 MB	6.410s	0.803s	5.392s	
ml_migratedb.ml_customer	0	30000	30000	16.2 MB	11.446s	1.169s	10.882s	
ml_migratedb.ml_tinyint1	0	201	201	4.3 kB	0.033s	0.002s	0.002s	
COPY Threads Completion	0	4	4		59.306s			
Reset Sequences	0	0	0		1.863s			
Create Foreign Keys	0	0	0		0.000s			
Install Comments	0	1	1		0.022s			
Total import time	✓	910201	910201	76.7 MB	1m1.191s			

Alternatively you can also check the platform to make sure all the tables are present by looking under the tables tab and making sure there are no active queries migration queries against the cluster (index creation, index backfill, copy from, create table, etc.)

Using `ysqlsh -h <IP> <dbname>` you can check out the tables that have been migrated.

```
ysqlsh (11.2-YB-2.7.2.0-b0)
Type "help" for help.

ml_migratedb=# \d
                                List of relations
 Schema | Name | Type | Owner
-----+-----+-----+-----
ml_migratedb | ml_customer | table | yugabyte
ml_migratedb | ml_order_line | table | yugabyte
ml_migratedb | ml_order_line_id_seq | sequence | yugabyte
ml_migratedb | ml_orders | table | yugabyte
ml_migratedb | ml_orders_id_seq | sequence | yugabyte
ml_migratedb | ml_tinyint1 | table | yugabyte
ml_migratedb | ml_tinyint1_id_seq | sequence | yugabyte
ml_migratedb | orders_orig | table | yugabyte
(8 rows)

ml_migratedb=#
```

Conclusion

The ability to migrate seamlessly to a new database helps to take the pressure off moving to a new technology. With the changes we've made to pgloader at Yugabyte to enable users to easily move your current MySQL database onto YugabyteDB, we relieve that headache. As always we are here to help and answer any questions you may have. Join us on our [community Slack](#) channel, and star us on [GitHub](#).